

```

# Importing necessary libraries for data manipulation, model building,
and evaluation
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt # Ensure this is imported for
plotting
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.neighbors import KNeighborsClassifier
from sklearn.neural_network import MLPClassifier
from sklearn.ensemble import RandomForestClassifier,
AdaBoostClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score, classification_report,
roc_auc_score
from sklearn.feature_selection import SelectKBest, f_classif
from imblearn.over_sampling import SMOTE

# Printing success message to confirm all libraries have been imported
without errors
print("Libraries imported successfully.")

```

Libraries imported successfully.

```

# Importing the pandas library, which provides data manipulation and
analysis functions
import pandas as pd

# Loading the dataset from the specified file path into a DataFrame
object called 'data'
data = pd.read_csv(r"C:\Users\teena\OneDrive\Desktop\healthcare-
dataset-stroke-data.csv")
# Printing a confirmation message indicating the dataset is loaded
successfully
print("Dataset loaded successfully.")
# Displaying the first few rows of the dataset using the head()
function
print(data.head())

```

Dataset loaded successfully.

	id	gender	age	hypertension	heart_disease	ever_married	\
0	9046	Male	67.0	0	1		Yes
1	51676	Female	61.0	0	0		Yes
2	31112	Male	80.0	0	1		Yes
3	60182	Female	49.0	0	0		Yes
4	1665	Female	79.0	1	0		Yes

work_type	Residence_type	avg_glucose_level	bmi
-----------	----------------	-------------------	-----

smoking_status \					
0	Private	Urban	228.69	36.6	formerly smoked
1	Self-employed	Rural	202.21	NaN	never smoked
2	Private	Rural	105.92	32.5	never smoked
3	Private	Urban	171.23	34.4	smokes
4	Self-employed	Rural	174.12	24.0	never smoked

	stroke
0	1
1	1
2	1
3	1
4	1

Selecting only the relevant features from the dataset for analysis
Keeping features that are relevant to predicting stroke, based on prior research or domain knowledge

```
data = data[['age', 'gender', 'hypertension', 'heart_disease', 'avg_glucose_level', 'bmi', 'smoking_status', 'stroke']]
```

Displaying a message to indicate the selection of relevant features has been done

```
print("Selected Relevant Features")
```

Displaying the first few rows of the selected data to verify the selected features

```
data.head()
```

Selected Relevant Features

	age	gender	hypertension	heart_disease	avg_glucose_level	bmi
\						
0	67.0	Male	0	1	228.69	36.6
1	61.0	Female	0	0	202.21	NaN
2	80.0	Male	0	1	105.92	32.5
3	49.0	Female	0	0	171.23	34.4
4	79.0	Female	1	0	174.12	24.0

	smoking_status	stroke
0	formerly smoked	1
1	never smoked	1

```

2      never smoked      1
3           smokes      1
4      never smoked      1

```

```

# Creating a copy of the original dataset to prevent modifying the
original data directly

```

```

data = data.copy()

```

```

# Filling missing values in the 'bmi' column with the mean of the
'bmi' column

```

```

# This is using the 'fillna()' method to replace NaN values with the
mean

```

```

# Setting inplace=False to avoid potential warnings by modifying the
data in a copied version

```

```

data['bmi'] = data['bmi'].fillna(data['bmi'].mean())

```

```

# Checking if missing values in the 'bmi' column have been
successfully handled

```

```

# Using 'isnull().sum()' to verify that there are no remaining NaN
values in 'bmi'

```

```

print("Missing values in 'bmi' column handled successfully.")

```

```

print(data['bmi'].isnull().sum())

```

```

Missing values in 'bmi' column handled successfully.

```

```

0

```

```

# Encoding the 'gender' column by mapping categorical values to
numeric values

```

```

# Mapping 'Female' to 0 and 'Male' to 1

```

```

data['gender'] = data['gender'].map({'Female': 0, 'Male': 1})

```

```

# Encoding the 'smoking_status' column by mapping categorical values
to numeric values

```

```

data['smoking_status'] = data['smoking_status'].map({
    'never smoked': 0,
    'smokes': 1,
    'formerly smoked': 2
})

```

```

# Printing a message to confirm that the encoding of 'gender' and
'smoking_status' is complete

```

```

print("Encoded Categorical Variables 'gender' and 'smoking_status'")

```

```

# Displaying the first few rows of the dataset to verify the changes
data.head()

```

```

Encoded Categorical Variables 'gender' and 'smoking_status'

```

```

   age  gender  hypertension  heart_disease  avg_glucose_level
bmi  \
0  67.0     1.0             0              1          228.69

```

36.600000					
1	61.0	0.0	0	0	202.21
28.893237					
2	80.0	1.0	0	1	105.92
32.500000					
3	49.0	0.0	0	0	171.23
34.400000					
4	79.0	0.0	1	0	174.12
24.000000					

	smoking_status	stroke
0	2.0	1
1	0.0	1
2	0.0	1
3	1.0	1
4	0.0	1

Defining features and target

```
X = data.drop('stroke', axis=1)
```

```
y = data['stroke']
```

```
print("Separated features and target variable.")
```

```
print(X.head())
```

```
print(y.head())
```

Separated features and target variable.

	age	gender	hypertension	heart_disease	avg_glucose_level
bmi \					
0	67.0	1.0	0	1	228.69
36.600000					
1	61.0	0.0	0	0	202.21
28.893237					
2	80.0	1.0	0	1	105.92
32.500000					
3	49.0	0.0	0	0	171.23
34.400000					
4	79.0	0.0	1	0	174.12
24.000000					

	smoking_status
0	2.0
1	0.0
2	0.0
3	1.0
4	0.0

0	1
1	1
2	1
3	1
4	1

Name: stroke, dtype: int64

```

# Importing train_test_split function from sklearn library to split
data into training and testing sets
from sklearn.model_selection import train_test_split

# Splitting the dataset into training and testing sets, where X
represents features and y represents target variable
# Setting test_size to 0.2, which means 20% of the data will be used
for testing, and 80% will be used for training
# Setting random_state to 42 to ensure reproducibility of the split;
using the same seed will give the same split every time
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# Printing a message indicating that the data has been split into
training and testing sets
print("Split Data into Training and Testing Sets")

# Printing the shape (number of samples and features) of the training
set to show its size
print("Training Set Size:", X_train.shape)

# Printing the shape (number of samples and features) of the testing
set to show its size
print("Testing Set Size:", X_test.shape)

Split Data into Training and Testing Sets
Training Set Size: (4088, 7)
Testing Set Size: (1022, 7)

from sklearn.preprocessing import StandardScaler
import numpy as np

# To Ensure X_train and X_test have been split before this point

# Step 1: Scaling the data
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train) # Scale training data
X_test_scaled = scaler.transform(X_test) # Scale testing data
print("Data scaled successfully.")

# Step 2: Checking for NaN values in scaled data
print("Checking for NaN values in scaled training data:")

# Summing up the count of NaN values for each column in X_train_scaled
# and storing it in 'nan_columns' to identify columns with NaN issues
nan_columns = np.isnan(X_train_scaled).sum(axis=0)
print("NaN counts per column:", nan_columns)

# Step 3: Replacing any NaN values with column means
if nan_columns.any(): # Check if there are any NaN values
    X_train_scaled = np.where(np.isnan(X_train_scaled),

```

```

np.nanmean(X_train_scaled, axis=0), X_train_scaled)
X_test_scaled = np.where(np.isnan(X_test_scaled),
np.nanmean(X_test_scaled, axis=0), X_test_scaled)
print("Replaced NaN values with column means.")
else:
    print("No NaN values found in scaled data.")

Data scaled successfully.
Checking for NaN values in scaled training data:
NaN counts per column: [ 0  0  0  0  0  0  0 1233]
Replaced NaN values with column means.

# Importing SMOTE from imbalanced-learn library to handle class
imbalance
smote = SMOTE(random_state=42) # Creating an instance of SMOTE with a
fixed random state for reproducibility

# Applying SMOTE to the training data to balance the class
distribution
X_train_resampled, y_train_resampled =
smote.fit_resample(X_train_scaled, y_train)

# Printing a message to confirm that SMOTE has been successfully
applied for class balancing
print("SMOTE applied for class balancing successfully.")


# Displaying the distribution of classes after SMOTE to verify
balancing
print("Class distribution after SMOTE:\n",
pd.Series(y_train_resampled).value_counts())

SMOTE applied for class balancing successfully.
Class distribution after SMOTE:
stroke
0    3901
1    3901
Name: count, dtype: int64

# Generating the correlation matrix
correlation_matrix = data.corr()

# Plotting the correlation matrix as a heatmap
plt.figure(figsize=(10, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm',
fmt=".2f", cbar=True)
plt.title("Correlation Matrix for Stroke Prediction Dataset")
plt.show()

```



```
from sklearn.feature_selection import SelectKBest, f_classif

# Setting the number of top features to select
# 'k' is set to the lesser of 16 or 20, to ensure it doesn't exceed
the number of features available
k = min(16, data.shape[1] - 1) # Adjusts k to the lesser of 16 or the
number of features

# Creating an instance of SelectKBest for feature selection
# SelectKBest is using the f_classif function to score features based
on their ANOVA F-value with the target
selector = SelectKBest(f_classif, k=k)

# Fitting the SelectKBest model on the resampled training data
X_train_selected = selector.fit_transform(X_train_resampled,
y_train_resampled)
```

```

# Transforming the scaled testing data to only include the selected
features
X_test_selected = selector.transform(X_test_scaled)

# Confirming the feature selection process
print("Top features selected successfully.")
# Printing the shape of the selected training features to verify the
expected number of features
print("Selected training features shape:", X_train_selected.shape)
# Printing the shape of the selected testing features to confirm
consistency with the training set
print("Selected testing features shape:", X_test_selected.shape)

```

```

Top features selected successfully.
Selected training features shape: (7802, 7)
Selected testing features shape: (1022, 7)

```

```

# Importing necessary libraries
from xgboost import XGBClassifier
from sklearn.linear_model import LogisticRegression

```

```

# Defining classifiers

```

```

classifiers = {
    'KNN': KNeighborsClassifier(),
    'MLP': MLPClassifier(max_iter=1000),
    'Random Forest': RandomForestClassifier(),
    'AdaBoost': AdaBoostClassifier(),
    'GaussianNB': GaussianNB(),
    'SVM (Linear)': SVC(kernel='linear', probability=True),
    'SVM (RBF)': SVC(kernel='rbf', probability=True),
    'SVM (Sigmoid)': SVC(kernel='sigmoid', probability=True),
    'XGBoost': XGBClassifier(use_label_encoder=False,
eval_metric='logloss'),
    'Logistic Regression': LogisticRegression(),
}

```

```

print("Classifiers defined successfully.")

```

```

Classifiers defined successfully.

```

```

from xgboost import XGBClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import AdaBoostClassifier,
RandomForestClassifier
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn.neural_network import MLPClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import accuracy_score, roc_auc_score,

```



```

classification_report
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# Defining classifiers with adjustments
classifiers = {
    'KNN': KNeighborsClassifier(),
    'MLP': MLPClassifier(max_iter=1000),
    'Random Forest': RandomForestClassifier(),
    'AdaBoost': AdaBoostClassifier(algorithm='SAMME'),
    'GaussianNB': GaussianNB(),
    'SVM (Linear)': SVC(kernel='linear', probability=True),
    'SVM (RBF)': SVC(kernel='rbf', probability=True),
    'SVM (Sigmoid)': SVC(kernel='sigmoid', probability=True),
    'XGBoost': XGBClassifier(eval_metric='logloss'),
    'Logistic Regression': LogisticRegression()
}

# Storing results
results = {
    "Classifier": [],
    "Accuracy": [],
    "AUC": [],
    "Precision": [],
    "Recall": [],
    "F1-score": []
}

# Training and evaluating classifiers
for name, clf in classifiers.items():
    clf.fit(X_train_selected, y_train_resampled)
    y_pred = clf.predict(X_test_selected)
    y_pred_proba = clf.predict_proba(X_test_selected)[:, 1] if
hasattr(clf, "predict_proba") else None

    accuracy = accuracy_score(y_test, y_pred)
    auc = roc_auc_score(y_test, y_pred_proba) if y_pred_proba is not
None else "N/A"
    report = classification_report(y_test, y_pred, output_dict=True)

    # Append results
    results["Classifier"].append(name)
    results["Accuracy"].append(accuracy)
    results["AUC"].append(auc)
    results["Precision"].append(report['weighted avg']['precision'])
    results["Recall"].append(report['weighted avg']['recall'])
    results["F1-score"].append(report['weighted avg']['f1-score'])

    print(f"\n{name} Classifier:")

```

```

print(f"Accuracy: {accuracy:.4f}")
print(f"AUC: {auc}")
print("Classification Report:")
print(classification_report(y_test, y_pred))

# Converting results to DataFrame
summary_df = pd.DataFrame(results)

# Displaying summary
print("\nSummary of Classifier Results:")
print(summary_df)

# Melting DataFrame for seaborn compatibility
summary_melted = summary_df.melt(id_vars="Classifier",
value_vars=["Accuracy", "Precision", "Recall", "F1-score"],
                                var_name="Metric",
                                value_name="Score")

# Plotting with seaborn
plt.figure(figsize=(12, 6))
sns.barplot(data=summary_melted, x="Classifier", y="Score",
hue="Metric")

# Adding plot labels and title
plt.title("Model Performance Metrics")
plt.xlabel("Classifier")
plt.ylabel("Score")
plt.ylim(0, 1) # Score range from 0 to 1 for better visibility
plt.legend(title="Metric")

# Displaying the plot
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()

```

KNN Classifier:

Accuracy: 0.8014

AUC: 0.6745631720430108

Classification Report:

	precision	recall	f1-score	support
0	0.95	0.83	0.89	960
1	0.12	0.37	0.18	62
accuracy			0.80	1022
macro avg	0.54	0.60	0.54	1022
weighted avg	0.90	0.80	0.84	1022

MLP Classifier:

Accuracy: 0.8288
AUC: 0.7763608870967742
Classification Report:

	precision	recall	f1-score	support
0	0.96	0.86	0.90	960
1	0.16	0.42	0.23	62
accuracy			0.83	1022
macro avg	0.56	0.64	0.57	1022
weighted avg	0.91	0.83	0.86	1022

Random Forest Classifier:
Accuracy: 0.8787
AUC: 0.7661542338709678
Classification Report:

	precision	recall	f1-score	support
0	0.95	0.92	0.93	960
1	0.17	0.26	0.21	62
accuracy			0.88	1022
macro avg	0.56	0.59	0.57	1022
weighted avg	0.90	0.88	0.89	1022

AdaBoost Classifier:
Accuracy: 0.6898
AUC: 0.8503864247311828
Classification Report:

	precision	recall	f1-score	support
0	0.98	0.68	0.80	960
1	0.14	0.84	0.25	62
accuracy			0.69	1022
macro avg	0.56	0.76	0.53	1022
weighted avg	0.93	0.69	0.77	1022

GaussianNB Classifier:
Accuracy: 0.7906
AUC: 0.833921370967742
Classification Report:

	precision	recall	f1-score	support
0	0.98	0.79	0.88	960
1	0.19	0.73	0.30	62
accuracy			0.79	1022

macro avg	0.58	0.76	0.59	1022
weighted avg	0.93	0.79	0.84	1022

SVM (Linear) Classifier:

Accuracy: 0.7407

AUC: 0.8562920026881721

Classification Report:

	precision	recall	f1-score	support
0	0.98	0.74	0.84	960
1	0.17	0.81	0.27	62

accuracy			0.74	1022
macro avg	0.57	0.77	0.56	1022
weighted avg	0.93	0.74	0.81	1022

SVM (RBF) Classifier:

Accuracy: 0.7583

AUC: 0.7935315860215054

Classification Report:

	precision	recall	f1-score	support
0	0.98	0.76	0.86	960
1	0.17	0.74	0.27	62

accuracy			0.76	1022
macro avg	0.57	0.75	0.56	1022
weighted avg	0.93	0.76	0.82	1022

SVM (Sigmoid) Classifier:

Accuracy: 0.6722

AUC: 0.7165658602150538

Classification Report:

	precision	recall	f1-score	support
0	0.97	0.67	0.79	960
1	0.12	0.69	0.20	62

accuracy			0.67	1022
macro avg	0.55	0.68	0.50	1022
weighted avg	0.92	0.67	0.76	1022

XGBoost Classifier:

Accuracy: 0.8581

AUC: 0.7767641129032258

Classification Report:

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	0.95	0.89	0.92	960
1	0.17	0.34	0.22	62
accuracy			0.86	1022
macro avg	0.56	0.62	0.57	1022
weighted avg	0.91	0.86	0.88	1022

Logistic Regression Classifier:

Accuracy: 0.7476

AUC: 0.8551579301075268

Classification Report:

	precision	recall	f1-score	support
0	0.98	0.75	0.85	960
1	0.16	0.77	0.27	62
accuracy			0.75	1022
macro avg	0.57	0.76	0.56	1022
weighted avg	0.93	0.75	0.81	1022

Summary of Classifier Results:

	Classifier	Accuracy	AUC	Precision	Recall	F1-score
0	KNN	0.801370	0.674563	0.902923	0.801370	0.844311
1	MLP	0.828767	0.776361	0.909435	0.828767	0.862762
2	Random Forest	0.878669	0.766154	0.903099	0.878669	0.890085
3	AdaBoost	0.689824	0.850386	0.933954	0.689824	0.770853
4	GaussianNB	0.790607	0.833921	0.930143	0.790607	0.841767
5	SVM (Linear)	0.740705	0.856292	0.933668	0.740705	0.807698
6	SVM (RBF)	0.758317	0.793532	0.929235	0.758317	0.819719
7	SVM (Sigmoid)	0.672211	0.716566	0.919682	0.672211	0.757841
8	XGBoost	0.858121	0.776764	0.906591	0.858121	0.879614
9	Logistic Regression	0.747554	0.855158	0.931292	0.747554	0.812385

